

Exploring Movie Recommendations with Apache Spark

Andres Rocha, Andrew Chan, Sophia Chung, Yohan Sofian

June 3, 2024

Abstract

This report details the development and functioning of a movie recommendation system built using Apache Spark. The system preprocesses data, computes feature vectors, and applies machine learning techniques to recommend movies based on textual and numerical data.

1 High-Level Overview

Introduction

The objective of this project is to build a scalable movie recommendation system using Apache Spark. This system leverages text processing and machine learning to suggest movies similar to a given query.

Setup

The project is implemented in Scala and uses Apache Spark, an open-source unified analytics engine for large-scale data processing. The environment is set up as follows:

```
val spark = SparkSession.builder()  
  .appName("MovieRecommendation")  
  .master("local[*]")  
  .getOrCreate()
```

```
Logger.getLogger("org").setLevel(Level.OFF)
Logger.getLogger("akka").setLevel(Level.OFF)
```

Data Loading and Preprocessing

Data is loaded from a CSV file containing top-rated movies. Each movie's 'overview' is preprocessed to replace nulls with a default string. Here, text preprocessing involves tokenization and removal of stop words to clean the data for further analysis.

Feature Engineering

The system creates feature vectors using several techniques:

- **TF-IDF:** Converts text data into a numeric form, reflecting the importance of words within the dataset.
- **Vector Assembler:** Combines features from different sources (TF-IDF vectors and numerical attributes like ratings and popularity) into a single feature vector.
- **Standard Scaler:** Standardizes features by scaling to unit variance.

Machine Learning Pipeline

A pipeline is constructed with various stages from tokenization to scaling. This pipeline automates the workflow of transforming and assembling data:

```
val pipeline = new Pipeline()
    .setStages(Array(tokenizer, remover, hashingTF, idf, assembler, scaler))
```

Recommendation Engine

The system uses cosine similarity to find movies that are most similar to a given query. This measure helps identify movies with similar feature vectors, hence likely to be of interest to the user.

Example Query and Results

An example query is processed through the system to find movies similar to "Spider-Man: Across the Spider-Verse", demonstrating the effectiveness of the system.

title	overview	vote_average	vote_count	popularity	runtime	similarity
Spider-Man: Across the Spider-Verse	After reuniting with director Robert Rodriguez, the Spidey team returns to save the multiverse.	8.8	1160	2859.047	140	1.0
Midnight Cowboy	"Joe Buck is a wide-eyed, naive young guy from Texas who is naive enough to think his life will change when he moves to New York City to become a hustler and find love. And though he does find love, he never really becomes the hustler he thought he would be. Instead, he becomes a hustler who finds love."	7.5	1151	17.437	127	0.29027100842634684
Giant Spider	In a mysterious land, a young boy befriends a giant spider that can talk and has the power to control other spiders.	7.2	57	18.372	84	0.19111307153145093
Eight Legged Freaks	The residents of a small town are terrorized by a giant, mutated spider.	5.7	1059	346.156	99	0.18455005708099004
The Amazing Spider-Man	For Peter Parker, being a superhero is just the added benefit of helping his friends and neighbors.	6.5	12057	163.998	141	0.17653226359680757

Conclusion

This movie recommendation system showcases the power of Apache Spark in handling and analyzing large datasets. Through effective preprocessing and feature engineering, it offers a robust platform for movie recommendations.

2 In-Depth Analysis of Code

Introduction

The provided Scala code implements a movie recommendation system using Apache Spark and its machine learning library (MLlib). This document provides an in-depth analysis of each part of the code to explain its functionality and purpose.

Imports

The code begins with several import statements:

```
1 import org.apache.log4j.{Level, Logger,
   ↪ PropertyConfigurator}
2 import org.apache.spark.sql.{DataFrame, SparkSession
   ↪ }
3 import org.apache.spark.ml.feature.{HashingTF, IDF,
   ↪ RegexTokenizer, StandardScaler,
   ↪ StopWordsRemover, VectorAssembler}
4 import org.apache.spark.ml.linalg.Vector
5 import org.apache.spark.ml.{Pipeline, PipelineModel}
6 import org.apache.spark.sql.functions._
7 import org.apache.spark.ml.linalg.Vectors
```

These imports bring in the necessary libraries for logging, Spark SQL, and MLlib features. Specifically, they include components for text processing, feature transformation, and vector operations.

Main Object and Method

The main object `MovieRecommendation` contains the entry point of the program:

```
1 object MovieRecommendation {  
2   def main(args: Array[String]): Unit = {  
3     val spark = SparkSession.builder()  
4       .appName("MovieRecommendation")  
5       .master("local[*]")  
6       .getOrCreate()
```

This initializes a Spark session named `MovieRecommendation` and sets it to run locally on all available cores.

Logging Configuration

The logging configuration is set to suppress unnecessary logs and configure logging properties:

```
1   Logger.getLogger("org").setLevel(Level.OFF)  
2   Logger.getLogger("akka").setLevel(Level.OFF)  
3   PropertyConfigurator.configure("/Users/  
    ↪ andresrocha/Downloads/CSC369/Lab6/src/main/  
    ↪ resources/log4j.properties")
```

This suppresses logs from `org` and `akka` packages and sets the logging configuration file.

Data Loading and Preprocessing

The dataset is loaded and preprocessed:

```
1  import spark.implicit._
2
3  val df = spark.read.option("header", "true")
4    .csv("top_1000_popular_movies_tmdb.csv")
5    .withColumn("vote_average", $"vote_average".
6      ↪ cast("double"))
7    .withColumn("vote_count", $"vote_count".cast("
8      ↪ double"))
9    .withColumn("popularity", $"popularity".cast("
10     ↪ double"))
11   .withColumn("runtime", $"runtime".cast("double
12     ↪ "))
```

The dataset is read from a CSV file, and specific columns are cast to double data types for further processing.

Handling Missing Data

Null or empty 'overview' values are replaced with a default value:

```
1  val defaultOverview = "No overview available"
2  val cleanedDF = df.withColumn("overview", when(
    ↳ col("overview").isNull || length(trim(col("
    ↳ overview"))) == 0, defaultOverview).
    ↳ otherwise(col("overview")))
```

Median values are calculated for numerical columns, and missing values are filled:

```
1  val voteAverageMedian = cleanedDF.stat.
    ↳ approxQuantile("vote_average", Array(0.5),
    ↳ 0.001).head
2  val voteCountMedian = cleanedDF.stat.
    ↳ approxQuantile("vote_count", Array(0.5),
    ↳ 0.001).head
3  val popularityMedian = cleanedDF.stat.
    ↳ approxQuantile("popularity", Array(0.5),
    ↳ 0.001).head
4  val runtimeMedian = cleanedDF.stat.
    ↳ approxQuantile("runtime", Array(0.5),
    ↳ 0.001).head
5
6  val filledDF = cleanedDF.na.fill(Map(
7    "vote_average" -> voteAverageMedian,
8    "vote_count" -> voteCountMedian,
9    "popularity" -> popularityMedian,
10   "runtime" -> runtimeMedian
11  ))
```

Text Preprocessing

The text in the 'overview' column is tokenized, stop words are removed, and features are hashed and transformed using TF-IDF:

```
1  val tokenizer = new RegexTokenizer()
2      .setInputCol("overview")
3      .setOutputCol("tokens")
4      .setPattern("\\W")
5
6  val remover = new StopWordsRemover()
7      .setInputCol("tokens")
8      .setOutputCol("filtered_tokens")
9
10 val hashingTF = new HashingTF()
11     .setInputCol("filtered_tokens")
12     .setOutputCol("raw_features")
13     .setNumFeatures(10000)
14
15 val idf = new IDF()
16     .setInputCol("raw_features")
17     .setOutputCol("tfidf_features")
```

Feature Combination and Scaling

Features are combined and scaled:

```
1  val assembler = new VectorAssembler()
2      .setInputCols(Array("tfidf_features", "
        ↪ vote_average", "vote_count", "popularity"
        ↪ , "runtime"))
3      .setOutputCol("features")
4      .setHandleInvalid("skip")
5
6  val scaler = new StandardScaler()
7      .setInputCol("features")
8      .setOutputCol("scaled_features")
9      .setWithStd(true)
10     .setWithMean(false)
```


Pipeline and Model Fitting

A pipeline is created to streamline the preprocessing steps, and the model is fitted:

```
1  val pipeline = new Pipeline().setStages(Array(  
    ↳ tokenizer, remover, hashingTF, idf,  
    ↳ assembler, scaler))  
2  
3  val model = pipeline.fit(filledDF)  
4  val processedDF = model.transform(filledDF)
```

Cosine Similarity Calculation

Cosine similarity is calculated between feature vectors:

```
1  def cosineSimilarity(v1: Vector, v2: Vector):  
    ↳ Double = {  
2      val dotProduct = v1.toArray.zip(v2.toArray).  
        ↳ map { case (x, y) => x * y }.sum  
3      val normA = math.sqrt(v1.toArray.map(x => x *  
        ↳ x).sum)  
4      val normB = math.sqrt(v2.toArray.map(x => x *  
        ↳ x).sum)  
5      dotProduct / (normA * normB)  
6  }
```

Finding Nearest Neighbors

The function `findNearestNeighbors` finds the most similar movies based on a query and numerical data:

```
1  def findNearestNeighbors(query: String,
   ↪ numericalData: Array[Double], k: Int = 5):
   ↪ DataFrame = {
2  val queryDF = Seq((query, numericalData(0),
   ↪ numericalData(1), numericalData(2),
   ↪ numericalData(3))).toDF("overview", "
   ↪ vote_average", "vote_count", "popularity"
   ↪ , "runtime")
3  val queryProcessedDF = model.transform(queryDF
   ↪ )
4  val queryFeatures = queryProcessedDF.select("
   ↪ scaled_features").first().getAs[Vector]("
   ↪ scaled_features")
5
6  val similarities = processedDF.select("title",
   ↪ "overview", "vote_average", "vote_count"
   ↪ , "popularity", "runtime", "
   ↪ scaled_features").as[(String, String,
   ↪ Double, Double, Double, Double, Vector)].
   ↪ map {
7  case (title, overview, voteAvg, voteCount,
   ↪ popularity, runtime, features) =>
8  val similarity = cosineSimilarity(
   ↪ queryFeatures, features)
9  (title, overview, voteAvg, voteCount,
   ↪ popularity, runtime, similarity)
10 }
11
12 val nearestNeighbors = similarities.sort($"_7"
   ↪ .desc).take(k)
13 spark.createDataFrame(nearestNeighbors).toDF("
   ↪ title", "overview", "vote_average", "
   ↪ vote_count", "popularity", "runtime", "
   ↪ similarity")
```

```
14 }
```

Example Query and Execution

An example query is executed to find similar movies:

```
1  val query = "After reuniting with Gwen Stacy,  
    ↳ Brooklyn's full-time, friendly neighborhood  
    ↳ Spider-Man is catapulted across the  
    ↳ Multiverse, where he encounters the Spider  
    ↳ Society, a team of Spider-People charged  
    ↳ with protecting the Multiverse's very  
    ↳ existence. But when the heroes clash on how  
    ↳ to handle a new threat, Miles finds  
    ↳ himself pitted against the other Spiders  
    ↳ and must set out on his own to save those  
    ↳ he loves most."  
2  val numericalData = Array(8.8, 1160, 2859.047,  
    ↳ 140)  
3  val nearestNeighbors = findNearestNeighbors(  
    ↳ query, numericalData)  
4  
5  nearestNeighbors.show()  
6  spark.stop()  
7  }  
8 }
```